

Module 3: Pipeline Approaches and Deterministic ER

Rebecca C. Steorts

Reading

- ▶ Binette and Steorts (2020)
- ▶ Christen (2012)

Agenda

- ▶ Pipeline Approach
- ▶ Deterministic Record Linkage
- ▶ Exact Matching
- ▶ Scoring Functions
- ▶ Application to El Salvador

Load R packages

```
library(RecordLinkage)
```

Data Cleaning Pipeline



Figure 1: Data cleaning pipeline.

Deterministic Record Linkage

The most commonly used record linkage methods are based on a series of deterministic rules involving the comparison of record attributes.

Exact Matching and off by k matching

- ▶ Exact matching is where two record pairs are linked if they agree on all common attributes.
- ▶ An extension, off by k -matching, states that two record pairs are a match if they match on all common attributes except at most k , where k is an integer larger than 0.
- ▶ Exact matching (or extensions) are used when all the attributes are categorical as it tends to perform well, as opposed to when textual variables are introduced.

Off by k matching example

Let's apply off by k matching to the RLdata500 data set, removing columns with missing values.

```
#load RLdata500  
library(blink)  
data(RLdata500)  
data <- RLdata500[-c(2,4)] # Remove  
head(data)
```

##	fname_c1	lname_c1	by	bm	bd
## 1	CARSTEN	MEIER	1949	7	22
## 2	GERD	BAUER	1968	7	27
## 3	ROBERT	HARTMANN	1930	4	30
## 4	STEFAN	WOLFF	1957	9	2
## 5	RALF	KRUEGER	1966	1	13
## 6	JUERGEN	FRANKE	1929	7	4

Off by k matching example

Now let's consider all possible pairs of records.

```
pairs = t(combn(1:nrow(RLdata500), 2))  
head(pairs)
```

```
##      [,1] [,2]  
## [1,]    1    2  
## [2,]    1    3  
## [3,]    1    4  
## [4,]    1    5  
## [5,]    1    6  
## [6,]    1    7
```

Off by k matching example

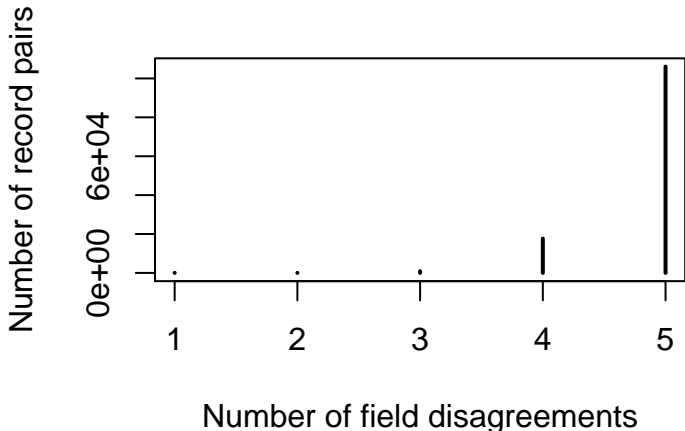
For each pair of record, compute the the number of fields which disagree. This takes a few minutes in R (there are more efficient ways to go about it).

```
n_disagree = sapply(1:nrow(pairs), function(i) {  
  recordA = data[pairs[i,1],]  
  recordB = data[pairs[i,2],]  
  sum(recordA != recordB)  
})
```

Off by k matching example

Now let's plot the distribution of the number of field disagreement among record pairs.

```
plot(table(n_disagree),  
      xlab="Number of field disagreements",  
      ylab="Number of record pairs")
```



Off by k matching example

We see most record pairs disagree on 4 or 5 fields. We can expect those pairs to be unmatched. Pairs disagreeing only on 1, 2 or 3 fields *might* be matches.

Off by k matching example

Exact matching: Link record pairs that for which there are no disagreement.

Here there are no such pairs...

```
sum(n_disagree == 0)
```

```
## [1] 0
```

Off by k matching example

Off by 1 matching: link record pairs which disagree only in one field.

```
# Links
```

```
links = pairs[n_disagree <= 1, ]
```

```
# Number of estimated links
```

```
nrow(links)
```

```
## [1] 46
```

```
# Number of correctly estimated links
```

```
sum(sapply(1:nrow(links), function(i) {  
  identity.RLdata500[links[i,1]] == identity.RLdata500[links[i,2]  
}))
```

```
## [1] 46
```

Off by k matching example

Off by 2 matching: link record pairs which disagree only in two field.

```
# Links
```

```
links = pairs[n_disagree <= 2, ]
```

```
# Number of estimated links
```

```
nrow(links)
```

```
## [1] 69
```

```
# Number of correctly estimated links
```

```
sum(sapply(1:nrow(links), function(i) {  
  identity.RLdata500[links[i,1]] == identity.RLdata500[links[i,2]  
}))
```

```
## [1] 50
```

Off by k matching example

Off by 3 matching: link record pairs which disagree only in two field.

```
# Links
```

```
links = pairs[n_disagree <= 3, ]
```

```
# Number of estimated links
```

```
nrow(links)
```

```
## [1] 975
```

```
# Number of correctly estimated links
```

```
sum(sapply(1:nrow(links), function(i) {  
  identity.RLdata500[links[i,1]] == identity.RLdata500[links[i,2]  
}))
```

```
## [1] 50
```


Scoring Rules

- ▶ Record attributes are often distorted by noise. Why would this occur?
- ▶ Linkage rules should account for such noise, distortions, and errors through scoring rules or functions.
- ▶ Examples commonly used for westernized names are the Edit (Levenshtein), Jaro, and Jaro-Winkler distance functions.

Edit (Levenshtein) distance (1966)

The Edit distance calculates the minimum number of substitutions required to transform a string s_1 into a string s_2 .

Formally,

$$\text{Edit} = 1 - \frac{L}{\text{maxLength}(s_1, s_2)}.$$

Example

Consider the number of substitutions required to transform from **Adam** to **Alan**. Use the Edit distance formulate to find the similarity score that is between $[0, 1]$.

Solution

The number of substitutions required is $L = 2$.

This is normalized into a similarity function using the following:

$$\text{Edit} = 1 - \frac{L}{\text{maxLength}(s_1, s_2)} = 1 - 2/4 = 1 - 0.5 = 0.5$$

Solution

Let's verify this in R.

```
s1 <- "Adam"  
s2 <- "Alan"  
levenshteinSim("s1", "s2")
```

```
## [1] 0.5
```

Jaro-Winkler

- ▶ The Jaro distance (1989), called J , considered common characters and character transpositions.
- ▶ The Jaro-Winkler (1990) similarity measure, denoted JW is:

$$JW(A, B) = J(A, B) + 0.1p(1 - J(A, B))$$

where p is the # of the first four characters that agree exactly.

Example

Let's return to the example of comparing Adam and Alan.

- ▶ Here, $p = 1$.
- ▶ Given the complexity, we will calculate J and JW using R.

Example

```
## It seems Jaro is not supported in R  
jarowinkler(s1,s2)
```

```
## [1] 0.7
```

e.g. Adam vs Alan: $p=1$, $J=0.67$ and $JW=0.7$.

These work well on English names that are less than 7 characters.

Other distance functions

There are many other distance functions, such as the Jaccard, Hamming, and Cosine distances just to name a few.

Case study on El Salvador.

Let's consider a case study from El Salvador that we will return to throughout the course.

Civil War in the Republic of El Salvador

Between 1980 and 1991, the Republic of El Salvador witnessed a civil war between the central government, the left-wing guerrilla Farabundo Marti National Liberation Front (FMLN), and right-wing paramilitary death squads.

After the peace agreement in 1992, the United Nations created a Commission on the Truth (UNTC) for El Salvador, which invited members of Salvadoran society to report war-related human rights violations, which mainly focused on killings and disappearances.

Civil War in the Republic of El Salvador

In order to collect such information the UNTC invited individuals through newspapers, radio, and television advertisements to come forward and testify.

The UNTC opened offices through El Salvador where witnesses could provide their testimonials, and this resulted in a list of potential victims with names, date of death, and reported location.

Data analysis

Let's read in the data and inspect it.

```
# read in data
df <- read.csv("./sv-mauricio/sv-mauricio.csv")
head(df)
```

##	X	ID	lastname	firstname	day	month	year	geocode	HandID	dept	muni
## 1	26	32	ASENSIO	ERNANDES	ALBERTO	NA	2 1981	150000	NA	15	NA
## 2	84	95	PALASIOS	AYALA	OBIDIO	NA	10 1985	150000	NA	15	NA
## 3	100	117		PALMA	SEBASTIAN	13	5 1980	40000	NA	4	NA
## 4	143	173		PERES	ARCADIO	NA	8 1984	40000	NA	4	NA
## 5	170	205	MAYA	QUESADA	ANTONIO	22	9 1984	0	NA	0	NA
## 6	189	227		MEJIA	ALFONSO	13	5 1980	40000	NA	4	NA

First steps

What steps do we need to take before running any algorithms or methods on this data?

Meet with your group and discuss. Be prepared to discuss ideas and any observations.

Task 1

What types of string distance metrics are appropriate and which are not appropriate?

Task 2

How does exact matching perform on this data set? What about off-by-one matching?

Task 3

How would you build a decision rule for matches/non-matches based upon scoring rules. What would your scoring rule be? Write this up.

Task 4

Implement your scoring rule from Task 4 and compare it to exact and off-by-one matching. What are your findings?

Task 5

Give insights into how you might be able to improve deterministic approaches moving forward if you re-did your analysis.