Module 4: Deterministic Blocking

Rebecca C. Steorts

joint with Olivier Binette

Reading

- Binette and Steorts (2020)
- Steorts, Ventura, Sadinle, Fienberg (2014)
- Murray (2016)
- Christen (2012), Chapter 4

Agenda

- Data Cleaning Pipeline
- Blocking
- Traditional Blocking
- Probabilistic Blocking

Load R packages and data

```r
knitr::opts_chunk$set(echo = TRUE,
                      fig.width=4,
                      fig.height=3,
                      fig.align="center")
library(RecordLinkage)
library(blink)
library(italy)
library(tidyverse)
library(assert)
data(italy08)
data(italy10)
data(RLdata500)
```

Data Cleaning Pipeline
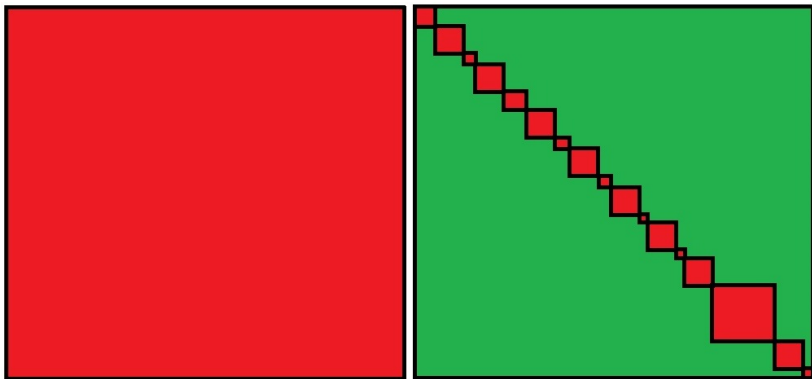

Figure 1: Data cleaning pipeline.

Blocking

Figure 2: Left: All to all record comparison. Right: Example of resulting blocking partitions.

Blocking

- Blocking partitions similar records into partitions/blocks.
- ER is only performed within each blocks.

Traditional Blocking

- A deterministic (fixed) partition is formed based upon the data.
- A partition is created by treating certain fields that are thought to be nearly error-free as fixed.
- Benefits: simple, easy to understand, and fast to implement.
- Downsides: the blocks are treated as error free, which is not usually accurate and can lead to errors in the ER task that cannot be accounted for.

Example: Blocking on date of birth year.

Probabilistic Blocking

▶ A probability model is used to cluster the data into blocks/partitions.

Example: Fellegi-Sunter (1969), or Locality Sensitive Hashing

Under both blocking approaches, record pairs that do not meet the blocking criteria are automatically classified as non-matches.
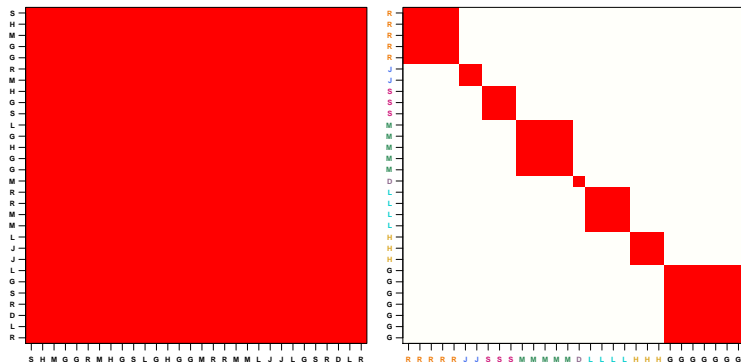
Example: Traditional blocking

Figure 3: All-to-all record comparisons (left) versus partitioning records into blocks by lastname initial and comparing records only within each partition (right).

Example: RLdata500

```
head(RLdata500)
```

```
##   fname_c1 fname_c2 lname_c1 lname_c2   by bm bd
## 1  CARSTEN     <NA>    MEIER     <NA> 1949  7 22
## 2     GERD     <NA>    BAUER     <NA> 1968  7 27
## 3   ROBERT     <NA> HARTMANN     <NA> 1930  4 30
## 4   STEFAN     <NA>    WOLFF     <NA> 1957  9  2
## 5     RALF     <NA>  KRUEGER     <NA> 1966  1 13
## 6  JUERGEN     <NA>   FRANKE     <NA> 1929  7  4
```

RLdata500 (Continued)

```r
# Total number of all to all record comparisons
choose(500,2)
```

```
## [1] 124750
```

RLdata500 (Continued)

```r
# Block by last name initial
last_init <- substr(RLdata500[,"lname_c1"], 1, 1)
head(last_init)
```

```
## [1] "M" "B" "H" "W" "K" "F"
```

```r
# Total number of blocks
length(unique(last_init))
```

```
## [1] 20
```

RLdata500 (Continued)

```r
# Total number of records per block
recordsPerBlock <- table(last_init)
head(recordsPerBlock)
```

```
## last_init
##  A  B  D  E  F  G
##  5 56  2  6 38 12
```

RLdata500 (Continued)

```r
# Block sizes can vary
plot(recordsPerBlock,
     cex.axis=0.6, xlab="", ylab="")
```

RLdata500 (Continued)

```r
# Total number of records pairs per block
choose(recordsPerBlock, 2)
```

```
## last_init
##    A    B    D    E    F    G    H    J    K    L    M
##   10 1540    1   15  703   66  496   28 1035   78 2850
##    T    V    W    Z
##    1   21 1326   10
```

```r
# Reduction on comparison space
sum(choose(recordsPerBlock, 2))
```

```
## [1] 14805
```

RLdata500 (Continued)

What is the overall dimension reduction form the original space to the reduced space induced by blocking?

Recall the original space of comparisons was

```
choose(500, 2)
```

[1] 124750

We have reduced the number of comparisons to

```
sum(choose(recordsPerBlock, 2))
```

[1] 14805

How do we calculate the reducation ratio?

The reduction ratio is

$$RR = \% \text{ comparisons eliminated by blocking.}$$

```
(choose(500, 2) - sum(choose(recordsPerBlock, 2))) /
choose(500, 2)
```

```
## [1] 0.8813226
```

How do we calculate the reducation ratio (via a function)?

```
reduction.ratio <- function(block.labels) {
  n_all_comp = choose(length(block.labels), 2)
  n_block_comp = sum(choose(table(block.labels), 2))

  (n_all_comp - n_block_comp) / n_all_comp
}

reduction.ratio(last_init)

## [1] 0.8813226
```

Pairwise Precision

```r
clusters = split(1:length(last_init), identity.RLdata500)

# Number of matching pairs among blocks
n_matches = sapply(clusters, function(records) {
  # Number of matches in that block
  sum(choose(table(identity.RLdata500[records]), 2))
})

# Total number of pairs
n_pairs = sum(choose(table(last_init), 2))

sum(n_matches) / n_pairs
```

```
## [1] 0.003377237
```

Pairwise Precision

```r
precision <- function(block.labels, IDs) {
  assert(length(block.labels) == length(IDs))

  clusters = split(1:length(block.labels), block.labels)

  # Number of matching pairs among blocks
  n_matches = sapply(clusters, function(records) {
    sum(choose(table(IDs[records]), 2))
  })

  # Total number of pairs
  n_pairs = sum(choose(table(block.labels), 2))

  sum(n_matches) / n_pairs
}

precision(last_init, identity.RLdata500)
```

[1] 0.003377327

Pairwise Recall

```
recall <- function(block.labels, IDs) {
  assert(length(block.labels) == length(IDs))

  precision(IDs, block.labels)
}
```

Italian Survey on Household and Wealth (SHIW)

We will now explore a case study to the SHIW

SHIW

- The Italian Survey on Household and Wealth (SHIW) is a sample survey 383 households conducted by the Bank of Italy every two years (2008 and 2010).
- The data set is anonymized to remove first and last name (and other sensitive information).

SHIW

The following attribute information is available:

- ▶ PARENT (parental status)
- ▶ GENDER
- ▶ ANASC (year of birth)
- ▶ NASCREG (working status)
- ▶ CIT (employment status)
- ▶ ACOM4C (branch of activity)
- ▶ STUDIO (town size)
- ▶ Q (quality of life status)
- ▶ QUAL (whether or not Italian national)
- ▶ SETT (highest educational level obtained)
- ▶ IREG (region of italy)

Explore Data

```r
head(italy08) # first year of SHIW
```

```
##          id PARENT SEX ANASC NASCREG CIT ACOM4C STUDIO Q QUAL SETT IREG
## 1 1040021      1   2  1948      16   1      0      5 1    2    3   16
## 2 1040022     10   2  1952      16   1      0      7 1    2    3   16
## 3 1110521      1   1  1972      20   1      2      5 1    1    4   20
## 4 1110522      3   1  1935      20   1      2      2 3    6    5   20
## 5 1110523      3   2  1941      20   1      2      3 3    6    5   20
## 6  119401      1   1  1941       7   1      0      4 3    6    5    7
```

Explore Data

```
head(italy08) # second year of SHIW
```

```
##          id PARENT SEX ANASC NASCREG CIT ACOM4C STUDIO Q QUAL SETT IREG
## 1 1040021      1   2  1948      16   1      0      5 1    2    3   16
## 2 1040022     10   2  1952      16   1      0      7 1    2    3   16
## 3 1110521      1   1  1972      20   1      2      5 1    1    4   20
## 4 1110522      3   1  1935      20   1      2      2 3    6    5   20
## 5 1110523      3   2  1941      20   1      2      3 3    6    5   20
## 6  119401      1   1  1941       7   1      0      4 3    6    5    7
```

Reformat Data

```
id08 <- italy08$id
id10 <- italy10$id
id <- c(italy08$id, italy10$id) # combine the id
italy08 <- italy08[-c(1)] # remove the id
italy10 <- italy10[-c(1)] # remove the id
italy <- rbind(italy08, italy10)
head(italy)
```

```
##   PARENT SEX ANASC NASCREG CIT ACOM4C STUDIO Q QUAL SETT IREG
## 1      1   2  1948      16   1      0      5 1    2    3   16
## 2     10   2  1952      16   1      0      7 1    2    3   16
## 3      1   1  1972      20   1      2      5 1    1    4   20
## 4      3   1  1935      20   1      2      2 3    6    5   20
## 5      3   2  1941      20   1      2      3 3    6    5   20
## 6      1   1  1941       7   1      0      4 3    6    5    7
```

Your turn

- Construct a blocking criterion for the SHIW data set.
- Provide code to construct the blocks
- Are your blocks well balanced?
- What is the reduction ratio?
- What is the pairwise recall and precision?
- Would you recommend your blocking criterion for an ER task? Why or why not.

Hint: You might consider blocking on gender, regions (in Italy), or combinations of these. What do you find?

# Your turn solution

Let's block on gender.

```
# block by gender
blockByGender <- italy$SEX
recordsPerBlock <- table(blockByGender)
head(recordsPerBlock)

## blockByGender
##     1     2
## 14442 12993
```

Your turn solution

The block sizes are similar. But note, they are still quite large.

```
# Checking block sizes
plot(recordsPerBlock,
cex.axis=0.6, xlab="", ylab="")
```

Your turn solution

```r
print(rr <- reduction.ratio(blockByGender))
```

```
## [1] 0.4986234
```

We have reduced the overall space by rougly 50 percent.

Your turn solution

```r
precision(blockByGender, id)
```

```
## [1] 3.599727e-05
```

```r
recall(blockByGender, id)
```

```
## [1] 0.9113109
```

This is not an optimal blocking criterion as ideally, we would want both the precision and recall to be close to 1.

Your turn solution

Let's block on the twenty regions in Italy.

```
blockRule <- (italy$IREG)
(recordsPerBlock <- table(blockRule))
```

```
## blockRule
##    1    2    3    4    5    6    7    8    9   10   11   12   13   14   15   16
## 2530  164 2875  645 1938  789 1050 2497 2203 1070 1235 1443  707  478 2329 1461
##   17   18   19   20
##  455  595 1678 1293
```

```
plot(recordsPerBlock,
cex.axis=0.6, xlab="", ylab="")
```
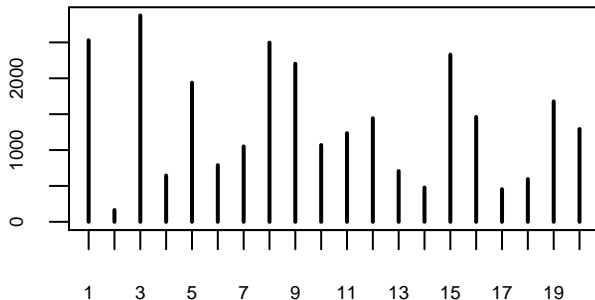


```
print(rr <- reduction.ratio(blockRule))
```

```
## [1] 0.9339148
```

```
precision(blockRule, id)
```

Your turn solution

Let's block on a combination of gender and region.

Hint: Use `tidyverse`.

The histogram is not printing now for some reason.

```
italy %>%
group_by(IREG, SEX) %>%
summarise(count = n(), .groups="drop") %>%
ggplot() +
geom_histogram(aes(count))
```

Your turn solution

```
blockIDs = paste(italy$IREG, italy$SEX, sep="_")
table(blockIDs)
```

```
## blockIDs
##  1_1  1_2 10_1 10_2 11_1 11_2 12_1 12_2 13_1 13_2 14_1 14_2 15_1 15_2 16_1 16_2
## 1282 1248  536  534  624  611  772  671  368  339  249  229 1310 1019  846  615
## 17_1 17_2 18_1 18_2 19_1 19_2  2_1  2_2 20_1 20_2  3_1  3_2  4_1  4_2  5_1  5_2
##  246  209  333  262  969  709   91   73  634  659 1489 1386  324  321 1056  882
##  6_1  6_2  7_1  7_2  8_1  8_2  9_1  9_2
##  418  371  517  533 1254 1243 1124 1079
print(rr <- reduction.ratio(blockIDs))
```

```
## [1] 0.9667954
precision(blockIDs, id)
```

```
## [1] 0.0005429847
recall(blockIDs, id)
```

```
## [1] 0.9103717
```