

Module X: fastlink

Rebecca C. Steorts

Reading

- ▶ Binette and Steorts (2020)
- ▶ Others ??

Probabilistic Entity Resolution

While Fellegi and Sunter (1969) have provided a framework for probabilistic entity resolution, there are few implementations that scale to large size data sets.

- ▶ Edmorando et al. (2020) developed fastlink a scalable implementation of the FS method.
- ▶ In addition, the authors incorporated auxiliary information such as population name frequency and migration rates.
- ▶ The authors used parallelization and hashing to merge millions of records in a near real-time on a laptop computer, and provided open-source software of their proposed methodology.

Agreement Patterns

- ▶ Two data sets (A and B) with variables in common
- ▶ Agreement value in field a for record pair (i, j)

$$\rho_a(i, j) = \begin{cases} \text{agree} \\ \text{disagree} \end{cases}$$

Agreement Patterns

	First	Last	Age	Street
Data set \mathcal{A}				
1	James	Smith	35	Devereux St.
Data set \mathcal{B}				
7	James	Smit	43	Dvereux St.

	agree	agree	disagree	agree

Agreement Patterns

	First	Last	Age	Street
Data set \mathcal{A}				
1	James	Smith	35	Devereux St.
Data set \mathcal{B}				
7	James	Smit	43	Dvereux St.
	agree	agree	disagree	agree

Agreement pattern $\gamma(i, j) = \{\gamma_1(i, j), \gamma_2(i, j), \dots, \gamma_K(i, j)\}$

Agreement Patterns

- ▶ We observe agreement patterns $\gamma(i, j)$
- ▶ We do not observe the matching status

$$C_{i,j} = \begin{cases} \text{non-match} \\ \text{match} \end{cases}$$

fastlink Model

$$\begin{aligned} C(i,j) &\stackrel{\text{iid}}{\sim} \text{Bernoulli}(\mu) \\ \rho(i,j) \mid C(i,j) = \text{non-match} &\stackrel{\text{iid}}{\sim} \mathcal{F}(\pi_{\text{NM}}) \\ \rho(i,j) \mid C(i,j) = \text{match} &\stackrel{\text{iid}}{\sim} \mathcal{F}(\pi_{\text{M}}) \end{aligned}$$

Where λ , π_{M} , π_{NM} are estimated via the EM algorithm

fastlink

- ▶ Available on CRAN
- ▶ We investigate it first on RLdata500

fastlink in R

```
library(fastLink)
library(RecordLinkage)
```

```
## Loading required package: DBI
```

```
## Loading required package: RSQLite
```

```
## Loading required package: ff
```

```
## Loading required package: bit
```

```
##
```

```
## Attaching package: 'bit'
```

```
## The following object is masked from 'package:base':
```

```
##
```

```
##      xor
```

```
## Attaching package ff
```

```
## - getOption("fftempdir")==="/var/folders/bv/xhclmwh90zg08"
```

```
## - getOption("fftempdir")=="/ff"
```

RLdata10000

```
# load RLdata10000
```

```
records <- read.table("data/RLdata10000.csv", sep="," , head  
head(records, 4)
```

##	fname_c1	fname_c2	lname_c1	lname_c2	by	bm	bd	rec_1
## 1	FRANK	<NA>	MUELLER	<NA>	1967	9	27	
## 2	MARTIN	<NA>	SCHWARZ	<NA>	1967	2	17	
## 3	HERBERT	<NA>	ZIMMERMANN	<NA>	1961	11	6	
## 4	HANS	<NA>	SCHMITT	<NA>	1945	8	14	

RLdata10000

```
# Number of unique records  
length(unique(records$ent_id))
```

```
## [1] 9000
```

```
#Linkage Fields
```

```
# linkage fields  
linkageFields <- c("fname_c1", "lname_c1", "by", "bm", "bd")
```

Exact Matching

```
# perform exact matching  
exact.match <- merge(records, records, by = linkageFields)
```

```
# number of self-matches  
sum(exact.match$rec_id.x == exact.match$rec_id.y)
```

```
## [1] 10000
```

```
# number of non-self matches  
sum(exact.match$rec_id.x != exact.match$rec_id.y)
```

```
## [1] 16
```

Who are they?

```
head(exact.match[exact.match$rec_id.x  
              != exact.match$rec_id.y,  
              c(linkageFields)], 4)
```

```
##          fname_c1 lname_c1    by bm bd  
## 973          BIRGIT ALBRECHT 1947 12  3  
## 974          BIRGIT ALBRECHT 1947 12  3  
## 1629 CHRISTINA    FRANKE 1997  8 13  
## 1630 CHRISTINA    FRANKE 1997  8 13
```

Preparation

```
# linkage fields
```

```
linkageFields <- c("fname_c1", "lname_c1", "by", "bm", "bd"
```

```
# string distance fields
```

```
stringDistFields <- c("fname_c1", "lname_c1")
```

```
# partial distance fields (fields where we allow  
# for agree, disagree, and partially agree)
```

```
partialMatchFields <- c("fname_c1", "lname_c1")
```


Run fastlink

```
out <- fastLink(dfA = records,
               dfB = records,
               varnames = linkageFields,
               stringdist.match = stringDistFields, # JW
               partial.match = partialMatchFields,
               cut.a = 0.94, cut.p = 0.84, # JW cutoffs
               dedupe = FALSE # 1-to-1 match
)
```

##

=====

fastLink(): Fast Probabilistic Record Linkage

=====

##

If you set return.all to FALSE, you will not be able to

dfA and dfB are identical, assuming deduplication of a s

Setting return.all to FALSE.

##

Calculating matches for each variable

fastlink Objects

fastLink has the following objects as output:

```
names(out)
```

```
## [1] "matches"    "EM"         "patterns"   "posterior"  "nob"
```

Who is matched

The indices of each matched pair can be found in `out$matches`

```
head(cbind(out$matches$inds.a, out$matches$inds.b), 6)
```

```
##      [,1] [,2]  
## [1,]    1    1  
## [2,]    2    2  
## [3,]    3    3  
## [4,]    4    4  
## [5,] 1957    4  
## [6,]    5    5
```

Counting Patterns

- ▶ Counts and FS weights for each patterns can be found in `outEMpatterns.w`
- ▶ Legend: 2 = Agree; 1 = Partially Agree; 0 = Disagree

```
tail(out$EM$patterns.w[, 1:7])
```

##		gamma.1	gamma.2	gamma.3	gamma.4	gamma.5	counts
##	[65,]	2	0	2	2	2	48 10
##	[66,]	0	1	2	2	2	20 3
##	[67,]	2	1	2	2	2	242 15
##	[68,]	0	2	2	2	2	56 8
##	[69,]	1	2	2	2	2	20 12
##	[70,]	2	2	2	2	2	10924 20

Matching Threshold

By default it is 0.85, but it can be easily changed:

```
out <- fastLink(dfA = records, dfB = records, varnames = 1:
stringdist.match = stringDistFields, partial.match = partial
cut.a = 0.94, cut.p = 0.84,
threshold.match = 0.90, # Matching threshold dedupe = FALSE
)
```

```
##
```

```
## =====
```

```
## fastLink(): Fast Probabilistic Record Linkage
```

```
## =====
```

```
##
```

```
## If you set return.all to FALSE, you will not be able to
```

```
## dfA and dfB are identical, assuming deduplication of a s
```

```
## Setting return.all to FALSE.
```

```
##
```

```
## Calculating matches for each variable.
```

```
## Getting counts for parameter estimation.
```

Exercises/TODO

- ▶ How would you calculate the precision and recall?
- ▶ Apply this to a much large data set with blocking. (Make this another module).
- ▶ Combine with fastlink and dblink