

Module X: fastlink

Rebecca C. Steorts

Reading

- ▶ Binette and Steorts (2020)
- ▶ Edmorando et al. (2020)
- ▶ Fellegi and Sunter (1969)

Probabilistic Entity Resolution

While Fellegi and Sunter (1969) have provided a framework for probabilistic entity resolution, there are few implementations that scale to large data sets.

Agenda

- ▶ We review fastlink, Edmorando et al. (2020)
- ▶ We illustrate a toy example on RLdata10000

- ▶ Edmorando et al. (2020) developed fastlink a scalable implementation of the FS method.
- ▶ In addition, the authors incorporated auxiliary information such as population name frequency and migration rates.
- ▶ The authors used parallelization and hashing to merge millions of records in a near real-time on a laptop computer, and provided open-source software of their proposed methodology.

Agreement Patterns

- ▶ Assume two data sets (A and B) with overlapping variables in common (such as name, gender, address, etc.)
- ▶ Define an agreement value in field a for record pair (i, j) :

$$\rho_a(i, j) = \begin{cases} \text{agree} \\ \text{disagree} \end{cases}$$

Agreement Patterns

	First	Last	Age	Street
Data set \mathcal{A}				
1	James	Smith	35	Devereux St.
Data set \mathcal{B}				
7	James	Smit	43	Dvereux St.
	agree	agree	disagree	agree

Agreement Patterns

	First	Last	Age	Street
Data set \mathcal{A}				
1	James	Smith	35	Devereux St.
Data set \mathcal{B}				
7	James	Smit	43	Dvereux St.
	agree	agree	disagree	agree

Agreement pattern $\gamma(i, j) = \{\gamma_1(i, j), \gamma_2(i, j), \dots, \gamma_K(i, j)\}$

Agreement Patterns

- ▶ We **observe** the agreement patterns $\gamma(i, j)$
- ▶ We **do not observe** the matching status

$$C_{i,j} = \begin{cases} \text{non-match} \\ \text{match} \end{cases}$$

fastLink Model

$$\begin{aligned} C(i,j) &\stackrel{\text{iid}}{\sim} \text{Bernoulli}(\mu) \\ \rho(i,j) \mid C(i,j) = \text{non-match} &\stackrel{\text{iid}}{\sim} \mathcal{F}(\pi_{\text{NM}}) \\ \rho(i,j) \mid C(i,j) = \text{match} &\stackrel{\text{iid}}{\sim} \mathcal{F}(\pi_{\text{M}}), \end{aligned}$$

where λ , π_{M} , π_{NM} are estimated via the EM algorithm

fastLink package

- ▶ Available on CRAN
- ▶ We investigate it first on RLdata10000
- ▶ We assume no blocking

fastlink in R

```
# load libraries  
library(fastLink)  
library(RecordLinkage)
```

RLdata10000

```
# load RLdata10000
records <-
  read.table("data/RLdata10000.csv",
             sep="," , header=TRUE)
head(records, 4)
```

##	fname_c1	fname_c2	lname_c1	lname_c2	by	bm	bd	rec_id	ent_id
## 1	FRANK	<NA>	MUELLER	<NA>	1967	9	27	1	3606
## 2	MARTIN	<NA>	SCHWARZ	<NA>	1967	2	17	2	2560
## 3	HERBERT	<NA>	ZIMMERMANN	<NA>	1961	11	6	3	3892
## 4	HANS	<NA>	SCHMITT	<NA>	1945	8	14	4	329

RLdata10000

```
# Number of unique records  
length(unique(records$ent_id))
```

```
## [1] 9000
```

Linkage Fields

```
# linkage fields  
linkageFields <- c("fname_c1", "lname_c1", "by", "bm", "bd")
```

Exact Matching

```
# perform exact matching
exact.match <- merge(records, records,
                     by = linkageFields)

# number of self-matches
sum(exact.match$rec_id.x == exact.match$rec_id.y)

## [1] 10000

# number of non-self matches
sum(exact.match$rec_id.x != exact.match$rec_id.y)

## [1] 16
```


Who are the non-self matches?

```
head(exact.match[exact.match$rec_id.x  
              != exact.match$rec_id.y,  
              c(linkageFields)], 4)
```

##	fname_c1	lname_c1	by	bm	bd
## 973	BIRGIT	ALBRECHT	1947	12	3
## 974	BIRGIT	ALBRECHT	1947	12	3
## 1629	CHRISTINA	FRANKE	1997	8	13
## 1630	CHRISTINA	FRANKE	1997	8	13

Preparation

```
# linkage fields
linkageFields <- c("fname_c1",
                  "lname_c1", "by", "bm", "bd")

# string distance fields
stringDistFields <- c("fname_c1", "lname_c1")

# partial distance fields (fields where we allow
# for agree, disagree, and partially agree)
partialMatchFields <- c("fname_c1", "lname_c1")
```

Run fastLink

```
out <- fastLink(dfA = records,
               dfB = records,
               varnames = linkageFields,
               # JW by default
               stringdist.match = stringDistFields,
               partial.match = partialMatchFields,
               # JW cutoffs
               cut.a = 0.94, cut.p = 0.84,
               # 1-to-1 match
               dedupe = FALSE)
```

##

=====

fastLink(): Fast Probabilistic Record Linkage

=====

##

If you set return.all to FALSE, you will not be able to

dfA and dfB are identical, assuming deduplication of a s

Setting return.all to FALSE

fastLink Objects

fastLink has the following objects as output:

```
names(out)
```

```
## [1] "matches"    "EM"         "patterns"   "posterior"  "nobs.a"
```

Who is matched?

The indices of each matched pair can be found in `out$matches`

```
head(cbind(out$matches$inds.a, out$matches$inds.b), 6)
```

```
##      [,1] [,2]  
## [1,]    1    1  
## [2,]    2    2  
## [3,]    3    3  
## [4,]    4    4  
## [5,] 1957    4  
## [6,]    5    5
```

Counting Patterns

- ▶ Counts and FS weights for each patterns can be found in `outEMpatterns.w`
- ▶ Legend: 2 = Agree; 1 = Partially Agree; 0 = Disagree

```
tail(out$EM$patterns.w[, 1:7])
```

##	gamma.1	gamma.2	gamma.3	gamma.4	gamma.5	counts	weig
## [65,]	2	0	2	2	2	48	10.493417296327
## [66,]	0	1	2	2	2	20	4.281720787087
## [67,]	2	1	2	2	2	242	14.922337282587
## [68,]	0	2	2	2	2	56	7.560387428981
## [69,]	1	2	2	2	2	20	12.604947813723
## [70,]	2	2	2	2	2	10924	18.201003924480

Matching Threshold

By default the **matching threshold 0.85**, but it can be easily changed:

```
out <- fastLink(dfA = records, dfB = records,
               varnames = linkageFields,
               stringdist.match = stringDistFields,
               partial.match = partialMatchFields,
               cut.a = 0.94, cut.p = 0.84,
               # Matching threshold dedupe = FALSE
               threshold.match = 0.90,)
```

##

=====

fastLink(): Fast Probabilistic Record Linkage

=====

##

If you set return.all to FALSE, you will not be able to

dfA and dfB are identical, assuming deduplication of a s

Setting return.all to FALSE.

Loading Packages

```
source("evaluationMetrics.R")  
#library(exchangeableER)  
library(magrittr)
```

```
##  
## Attaching package: 'magrittr'  
  
## The following object is masked from 'package:stringdist':  
##  
##      extract  
  
## The following object is masked from 'package:ff':  
##  
##      add
```

```
library(data.table)
```

```
##  
## Attaching package: 'data.table'
```

```
## The following object is masked from 'package:bit':
```


Setup

```
trueMembership <- records$ent_id
recordIds <- records$rec_id
numRecords <- dim(records)[1]
matches <-
  data.table(cbind(out$matches$inds.a,
                    out$matches$inds.b))
head(matches)
```

```
##      V1 V2
## 1:    1  1
## 2:    2  2
## 3:    3  3
## 4:    4  4
## 5: 1957  4
## 6:    5  5
```

```
dim(matches)[1]
```

```
## [1] 12072
```

True Positives, False Positives, and False Negatives

```
## True Positives, False Positives, and False Negatives:  
TP <- sum(records$ent_id[matches$V1]  
          == records$ent_id[matches$V2])  
FP <- sum(records$ent_id[matches$V1]  
          != records$ent_id[matches$V2])  
FN <- dim(matches)[1] - TP
```

FDR and FNR

```
## False Discovery Rate
```

```
FDR <- round(FP/(FP + TP), 2)
```

```
FDR
```

```
## [1] 0.01
```

```
## False Negative Rate
```

```
FNR <- round(FN/dim(matches)[1], 2)
```

```
FNR
```

```
## [1] 0.01
```

Precision and Recall

```
## Precision
precision <- round(1 - FDR, 2)

## Recall
recall <- round(1 - FNR, 2)

f1 <- round((2.0*TP)/(2.0*TP+FP+FNR), 2)

cbind(precision, recall, f1)

##      precision recall f1
## [1,]      0.99    0.99  1
```

TODO: The precision and recall is being done a bit weird here.

Exercises/TODO

- ▶ How would you calculate the precision and recall?
- ▶ Apply this to a much large data set with blocking. (Make this another module).
- ▶ Combine with fastlink and dblink